

Compression/Linear Prediction

Do You Want to Save on Disks?

- hopefully yes!
- Compression is the art of turning some data into a new set of “data” that is on average smaller.
- Choice of lossy vs. lossless.

Shannon Entropy

- Say I have m zeros, $(n-m)$ ones. How many ways are there to write?
- have n choose m possible orderings.
- Stirling: $n! \sim (n/e)^n$. n choose $m = (n/e)^n / ((n-m)/e)^{(n-m)} (m/e)^m = 1/(1-m/n)^{(n-m)} / (m/n)^m$.
- Take \log_2 : $\log_2() / n = -p_1 \log_2(p_1) - p_0 \log_2(p_0)$ where $p_1 = (n-m)/n$ and $p_0 = m/n$.
- Extension to many variables: $\log_2() / n = -\sum p_i \log_2(p_i)$.

Use in Compression

- The \log_2 of the number of orderings is just the original number times that sum, the Shannon Entropy (s).
- I can order all the possible strings with those probabilities, then tell you which one with a binary number of length ns.
- If entropy is small, ns will be shorter than original bit stream, and I can compress my data.
- Compression is lossless - I get back exactly the bits I started with.
- Example: 10% 1, 90%0: $s = -0.1 \cdot \log_2(0.1) - 0.9 \cdot \log_2(0.9) = .47$, so I should be able to save just over factor of 2 vs. original.

Entropy of Alphabets

- We can take wikipedia table of letter frequencies. If letters were randomly ordered, we'd expect to be able to compress.
- Of course, order is not random, so we can do much better. Shannon estimated 1 bit/letter.

```
#!/usr/bin/perl
use strict;
use warnings;
use Data::Dumper;

my %entropies = (
    'English' => 4.175759791063624,
    'French'   => 4.180328331429144,
    'German'   => 4.19151401615753,
    'Spanish'  => 4.211911779081105,
    'Portuguese' => 4.198258508748776,
    'Esperanto' => 4.153084400206026,
    'Italian'  => 4.060813730460597,
    'Turkish'  => 4.380740242124834,
    'Swedish'  => 4.306484306705044,
    'Polish'   => 4.558909447097535,
    'Dutch'    => 4.073673578455074,
    'Danish'   => 4.213486215928161,
    'Icelandic' => 4.502327036027735,
    'Finnish'  => 4.066543548089884,
    'Czech'    => 4.106298694824828,
    'Hawaiian' => 3.3423090035454903
);

print Dumper \%entropies;
```

How do we get there?

- Standard example is Huffman coding
- Make table of unique bit string for every unique symbol in original message
- To do: take two least-common symbols, assign 0/1 to them. merge symbols. Take next two least-common symbols, assign 1/0.
- If this one is merged, say assign merged to 1. Then assignments are 0 (for new symbol), 10 and 11 for old ones.
- Continue building tree. At end, have unique byte string for each symbol. On average, shorter than original.

Huffman in Practice

- Make a tree based on symbol probabilities, get unique mapping.
- Include special symbol for “we’re done”.
- Write byte string for each symbol, write “we’re done” at end.
- Reading: read tree, then read byte string until it matches tree entry. Gives first symbol. Repeat until “we’re done”.
- Effectively, we round p_i to nearest power of $1/2$. Will not be optimal, but usually not too bad.

Arithmetic Coding

- Is rounding of n_{bits} making you sad w/ Huffman?
- Let's map every message to a (very high precision!) number.
- Say 'a' has 10% probability. Then if number is between 0 and 0.1, we know first symbol is 'a'. Then multiply number by 10, repeat.
- End up with number that contains all information about original string, but with arbitrarily accurate entropies.
- There's an art to doing this efficiently...

Use Canned!

- Writing good low-level compression is hard!
- Our main goal will be to recast our data so that previous technique(s) will work well.
- Gzip2 does OK, bzip2 usually does better (but slower). Both have Huffman as part.
- If you understand what they're doing, you can set up your data in a way standard tools will work. Try a few, see which ones work the best for you.

Lossy vs. Lossless

- Say we have double precision data. It is *highly* unlikely all of those bits contain useful information.
- Say data have $\sigma=1$. I can round data to nearest $1/2^{\text{nbit}}$. Say I wanted to keep 4 bits of noise, then I would round my data to nearest $1/16$.
- This loses information. How much?
- Think of rounding as adding a noisy signal that brings data to nearest round value. How much noise does that add?
- Round to nearest 16th means max move is $1/32$, and average move is $1/64$. So, adjustment stream has $\sim 1/64^2$ times original variance.
- By rounding to nearest $1/16$, we have effectively increased variance by part in 64^2 . Amounts to throwing away less than 1 second per hour.
- 8 bits of noise: increased variance by $1/1024^2$, or 30 seconds per year. You can probably afford this...
- Generally, I can't imagine any typical physics situation where you need much more than 6 bits/number.

Modelling

- We can often win (big?!) by using previous data to guess at next data. A lot less information in that error, than in full value.
- LIGO example. Say we round data so there are 10 bits from $-\log(p)$.
- We know neighbors are similar, so if we save difference of point with previous, then $-\log(p)$ goes to 7.
- We've done a better job reducing entropy of our data, and as reward, should save 30% on storage space.
- Can we do better job predicting?

Linear Prediction

- Linear prediction is a powerful tool to interpolate/extrapolate based on a covariance.
- Say I have noise matrix, but am missing last data point.
- What value would I give it to minimize χ^2 ? $\chi^2 = \mathbf{d}^T \mathbf{N}^{-1} \mathbf{d}$.
- Differentiate w.r.t. last data point: $(0, 0, 0, \dots, 1)^T \mathbf{N}^{-1} \mathbf{d} = 0$.
- First vector kills everything but last row of \mathbf{N}^{-1} .
- $\text{sum}(\mathbf{N}^{-1}[-1, :-1] * \mathbf{d}(:n-1)) + \mathbf{N}^{-1}[-1, -1] * \mathbf{d}[-1] = 0$
- Let $\mathbf{v} = \mathbf{N}^{-1}[-1, :-1] / \mathbf{N}^{-1}[-1, -1]$, then $\mathbf{d}[-1] = \mathbf{v}^T \mathbf{d}[:-1]$

LP2


- This was used in South Africa to search for new gold mines, where N was the correlation between gold content as function of distance. Look up “Kriging”
- For 1D signals w/ stationary noise, this looks an awful lot like a convolution.
- We get to pick how many points to use in N^{-1} - for large datasets, you probably want to experiment.
- You may also want to add a (large) offset to N^{-1} . Essentially says you don't know what the mean is, so don't push answers towards zero.

LIGO In Practice

- Round raw data (at 64 bits) to some accuracy.
- Find N by taking normalized cross-correlation of LIGO with itself.
- Decide on # of points to use in prediction (~10 seems to work well). Find LP coefficients by inverting 10×10 block of N .
- From 10th point on, predict value given previous 9, save the *difference* between value and prediction.
- To restore, predict value given previous 9, and that to saved value, move one spot to right.
- Goes from 10.3 (raw) to 7.0 (diff) to 3.1 (LP). Bough more than a factor of 2!
- Final expected size is 51 kB, bzip2 gets 63 kB. Compare to 1 MB, we've saved factor of 20 with minimal loss.

LP Model

- Blue/yellow stars=true LIGO data
- Green is LP prediction from quantized data.
- We want to save difference between green & yellow.
- Note how green much closer to yellow than yellow neighbors to each other.



Practical Considerations

- Noise is often not quite stationary. In practice, this means you may want to keep # of points small.
- Often, a diff gets within 1/2 byte of optimal. Maybe that's good enough. Not the case for LIGO because of strong lines.
- Good LP has random errors (if they weren't, we could make a better model). Huffman/arithmetic encoding are well matched to this.
- Everything after digitizing/rounding is lossless. Whatever works well is what you should use!
- If you have complicated interpolation/extrapolation, I encourage more investigation into LP. It's often presented opaquely, but if you go back to χ^2 , you'll be fine.
- Look at `simple_compress_ligo.py` for fully worked example.